

Traffic Simulation in NYC



Team Violet Noise
Ameya, Emma, Helene, Kristin

Big Data Application Development CSCI GA 2437

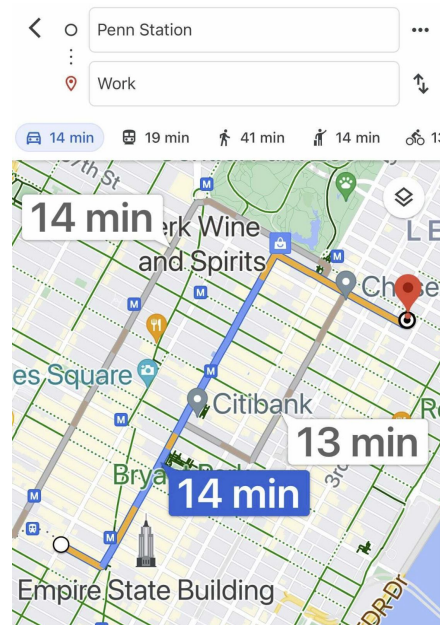

Motivation and Concept

Goals

1. Combine several main components of Apache Spark
 - Data processing
 - Machine learning
 - Graph algorithms
2. Try different types of algorithms to see how they handle graph features, and graph based noise

Concept

1. Simulate data for traffic in NYC
2. Assume that we collect the data through some IoT framework - we only have GPS type signals and general metadata about the graph
3. Attempt to solve for shortest path recommendation and vehicle identification



Pipeline



Data
Generation



Noise
Generation



Feature
Engineering

Model
Training

Vehicle
Classification

Vehicle Classification Pipeline

Feature
Engineering

Model
Training

Edge Weight
Forecast


Shortest Paths

Edge Weight Regression Pipeline

Data Generation

Data Generation


1. Created a graph to mimic midtown manhattan
 - Nodes are intersections
 - Edges are streets
 - “Source” and “Sink” nodes along the edge allow traffic flow
 - Edges have varying sizes, 2-4 lanes
2. Calculated all reasonable paths for every pair of nodes
3. Assign cars to these paths at regular intervals
4. Create instructions in XML for sumo to implement our traffic scenario
5. Extract trace data for all vehicles from sumo ~ **550 GB for 70 days**



| | | |
|---------|---------|---|
| 1.5 G | 4.4 G | /user/jl11257/big_data_project/traces/demo |
| 54.9 G | 164.8 G | /user/jl11257/big_data_project/traces/noised |
| 61.4 G | 184.2 G | /user/jl11257/big_data_project/traces/processed |
| 550.3 G | 1.6 T | /user/jl11257/big_data_project/traces/raw |

Features Embedded in Data Generation

1. Created time series for car arrival rates for each “source”
2. Different vehicles have different sizes and speed properties
3. Buses travel on a separate schedule, with specific routes and timed stops



```
1 cars = int(np.random.poisson(rate))
2 spread_cars = [min(i,2) for i in np.histogram(list(range(cars)),bins=900,density=False)[0]]
3 timestamps = [i + minute*60 + start_tm_seconds for i in range(900)]
4 cars_per_sec = [i for i in zip(timestamps, spread_cars) if i[1] > 0]
```

```
1 <vType accel="1.0" decel="3.5" id="Bus" length="15.0" maxSpeed="8.0" sigma="0.5" >
2 <vType accel="2.7" decel="4.6" id="Car1" length="4.0" maxSpeed="11.2" sigma="0.5">
3 <vType accel="2.4" decel="4.5" id="Car2" length="5.0" maxSpeed="11.2" sigma="0.5">
4 <vType accel="1.9" decel="4.3" id="Car3" length="7.0" maxSpeed="11.2" sigma="0.5">
```

```
1 {"busroute": "230|27",
2  "active_hrs": [5, 7, 9, 11, 13, 15, 17, 19, 21],
3  "start_min": 2,
4  "stops": ["station15"]}
```

Adding Noise

1. Random traffic jams

- 1-4 per day
- 3-9 min per jam

2. Making data more realistic

(~messy IOT device data)

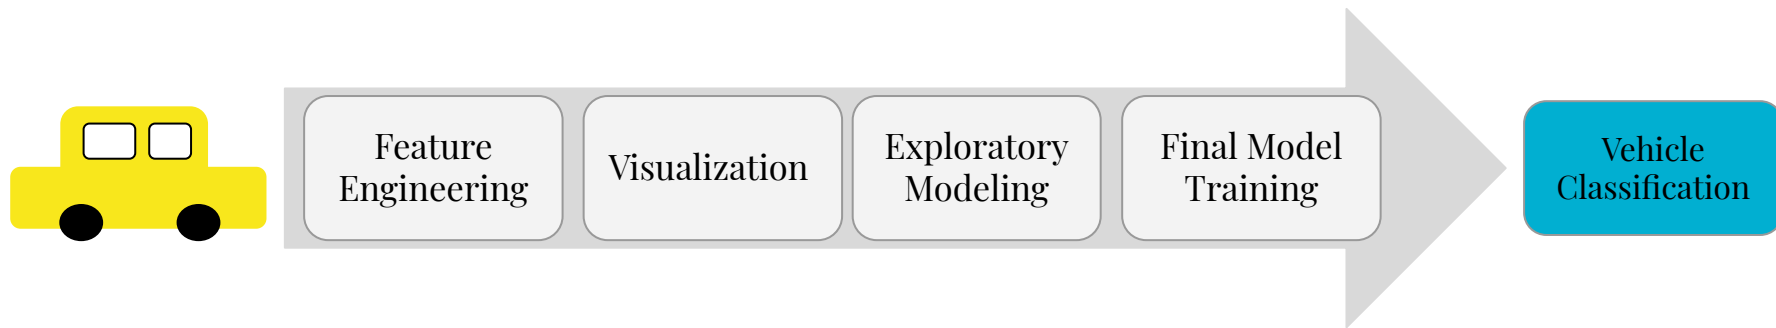
- 5% data loss from faulty GPS signals
- Cartesian coordinates ➡ GPS coordinates
- Dropped speed feature from raw data

```
1 val accident = udf((choice: Int) => {
2     val windowlength = if ((choice % 2) == 0) 3 else 5
3     val startTime = if ((choice % 2) == 0) 7 else 14
4     val duration = rand_gen.nextInt(max_accident_duration-min_accident_duration)+min_accident_duration
5     val start_time = rand_gen.nextInt(windowlength*60*60-duration+1)+(startTime*60*60)
6     val end_time = start_time+duration
7     val node = accident_candidates(rand_gen.nextInt(accident_candidates.length))(0).toString
8     Accident(start_time, end_time, node)
9 }
10 )
```

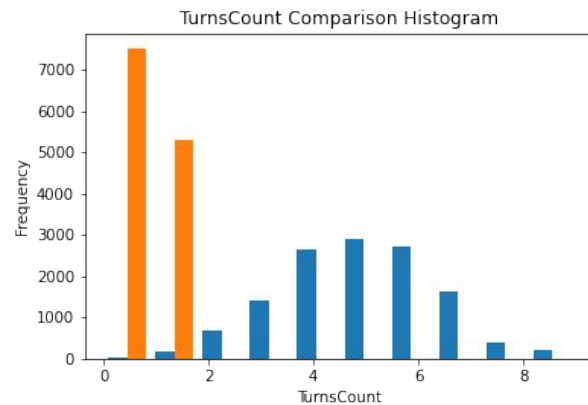
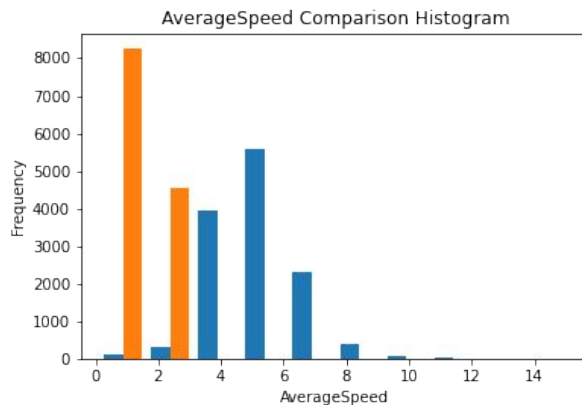
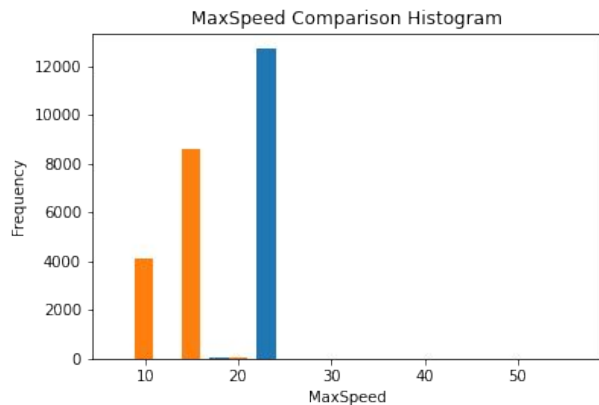
```
1 val transformXY = udf((x: Double, y: Double) => {
2     val xorig = if (x < 0) 0.0 else x / 100.0
3     val yorig = if (y < 0) 0.0 else y / 100.0
4     val mcoords = new DenseMatrix(3, 1, Array(xorig, yorig, 1))
5     val output = transformer.multiply(mcoords)
6     val latitude = BigDecimal(output.apply(0,0) / output.apply(2,0)).setScale(6,
7     BigDecimal.RoundingMode.HALF_UP).toDouble
8     val longitude = BigDecimal(output.apply(1,0) / output.apply(2,0)).setScale(6,
9     BigDecimal.RoundingMode.HALF_UP).toDouble
10    GpsCoord(latitude, longitude)
11 })
```

Vehicle Classification

Vehicle Classification Pipeline

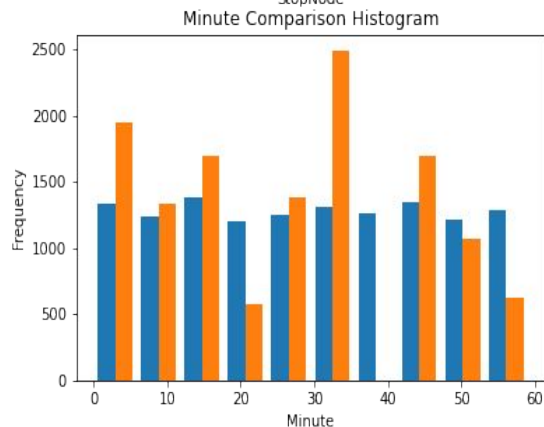
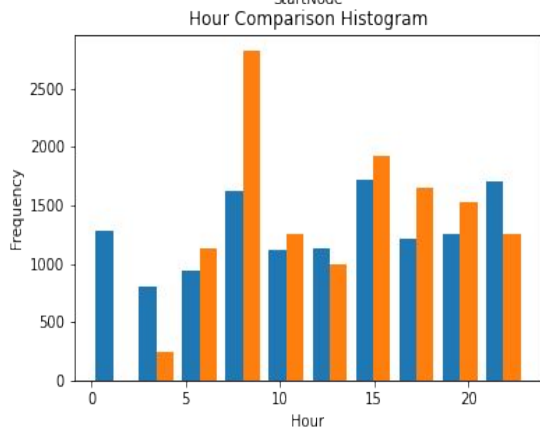

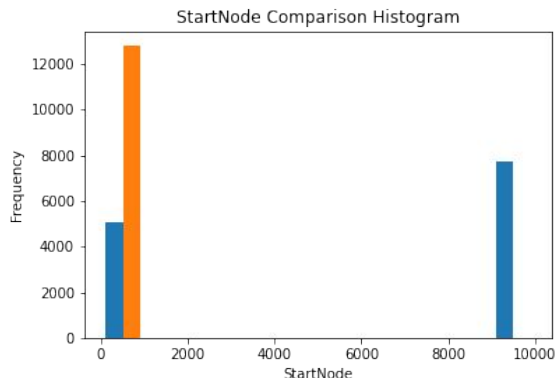


Feature Visualizations



MaxSpeed, AverageSpeed, and TurnsCount features of Car(blue) and Bus(orange) comparison

Feature Visualizations



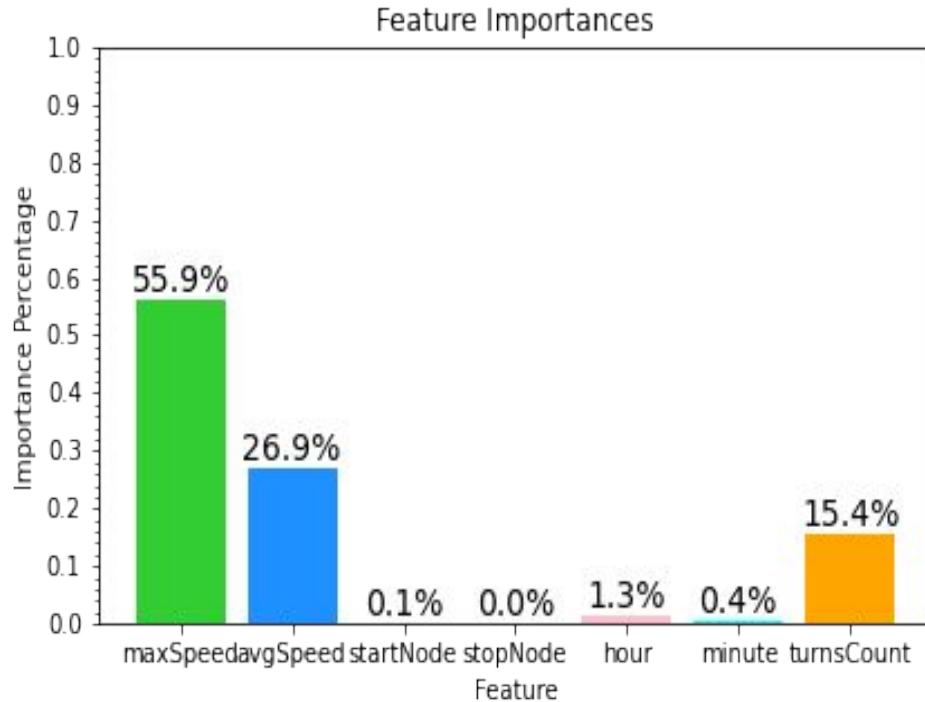
StartNode, StopNode, Hour, Minute features of Car(blue) and Bus(orange) comparison

Exploratory Model Results - Without Noise

| Model | auPR(Down Sampled Data) | auPR(Raw Test Data) |
|--------------------------|-------------------------|---------------------|
| Linear SVM | 0.010614 | 0.001081 |
| Linear Regression | 0.999609 | 0.998593 |
| Random Forest Classifier | 1.0 | 1.0 |

- Random Forest Classifier Model works best in these 3 models.
- We will focus on seeing how Random Forest Classifier Model reacts to noise in our data in the next stage.
- We also add cross validation to do parameter tuning and avoid overfitting.

Feature Importances Visualization



- MaxSpeed, AverageSpeed and TurnsCount were relevant features
- StartNode, StopNode, Hour and Minute were not that relevant features
- Relevant features were chosen for later model training

Cross Validation Results With Noise

| Model | auPR(Down Sampled Data) | auPR(Raw Test Data) | Features Included |
|--|-------------------------|---------------------|---|
| Random Forest Classifier combined with Cross Validator | 0.991220 | 0.919165 | 3 features (MaxSpeed, AverageSpeed and TurnsCount) |
| | 0.021272 | 0.002245 | 5 features (adding startNode and stopNode features) |
| | 0.992295 | 0.923198 | 4 features (adding hour feature) |

Summary of Findings

Adding startNode and stopNode features:

- The result gets worse.
- So not adding these two features to model training.

Adding Hour feature:

- The result improves a bit.
- So final model trained with four features: MaxSpeed, AverageSpeed, TurnsCount, Hour.

Sample Prediction Output from Test Data

There were 1474 buses and 655148 cars in the data set provided.

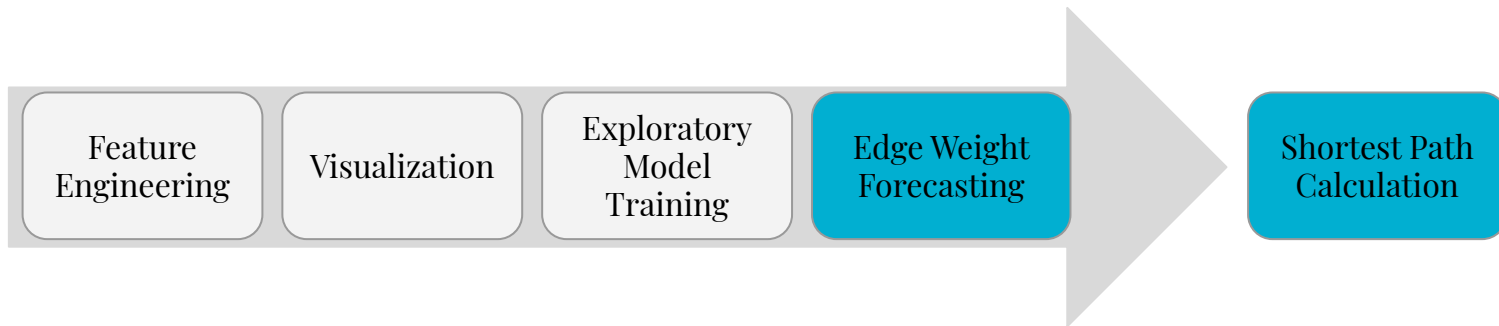
The model predicted 1595 buses and 655027 cars.

Area under precision-recall curve = 0.923198

Because we used auPR metric to train and the data is unbalanced, we will have slightly more false positives but with a better recall on the buses.

Edge Weight Forecasting

Shortest Path Pipeline



Feature Engineering

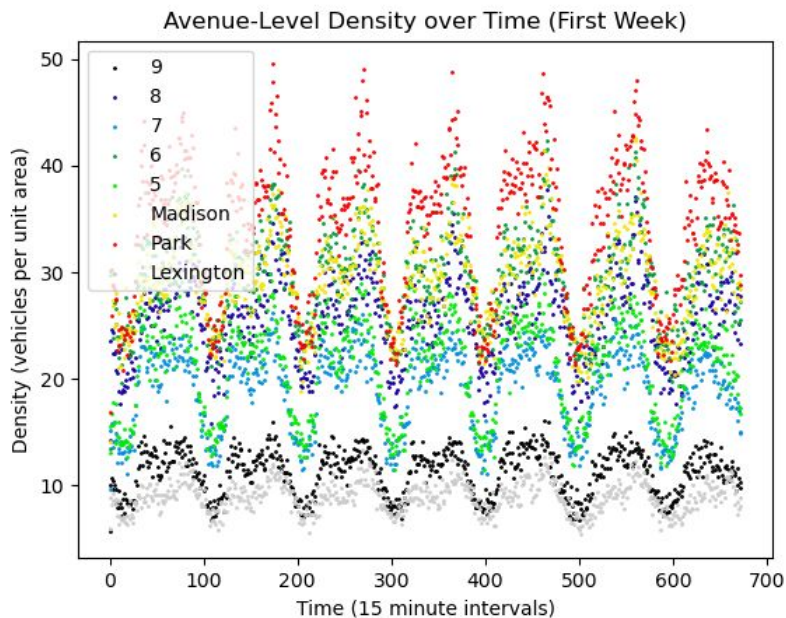
Obtained edge weight features using Spark SQL

- Customized measure of edge weight (for shortest path algorithm)
 - density: $\frac{\text{numOf Vehicles} \cdot \text{vehicleSize}}{\text{totalEdgeArea}}$
- Edge Features
 - density at time t : d_t
 - change of density: $\Delta d_{t-1,t}, \Delta d_{t-2,t}, \Delta d_{t-3,t}$
 - Time of the day as cyclical variable, formula we adapt $f(t) = \sin\left(\frac{\text{minuteOf Day} \cdot 2\pi}{24 \cdot 60}\right)$
- Graph Features
 - is edge two way & is edge in bus route
 - One hot encoding these two features

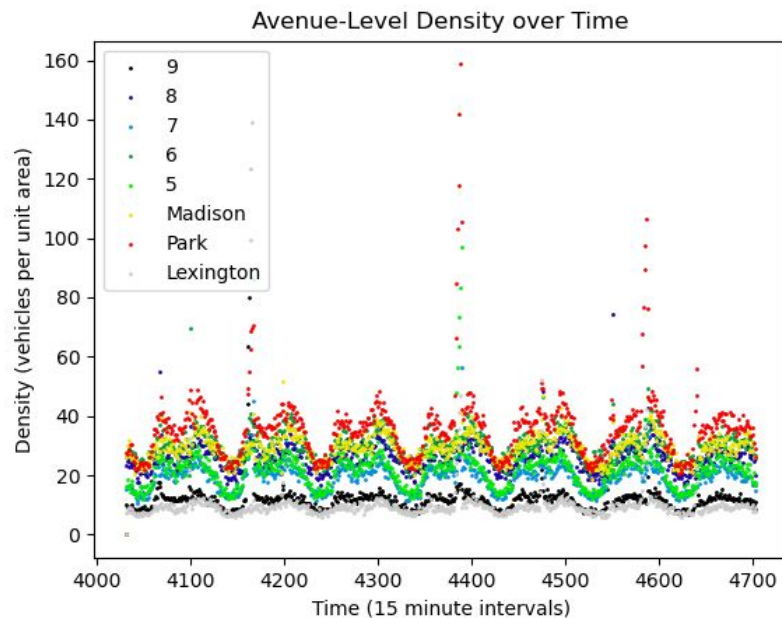
Custom train / test splits to handle time series data and data leakage

Feature Visualizations

Before Noise (1 week)




After Noise (1 week)



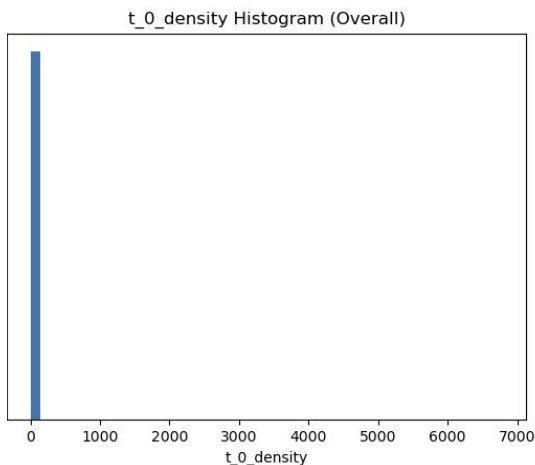
Feature Visualizations

Distributions of the Independent Variable

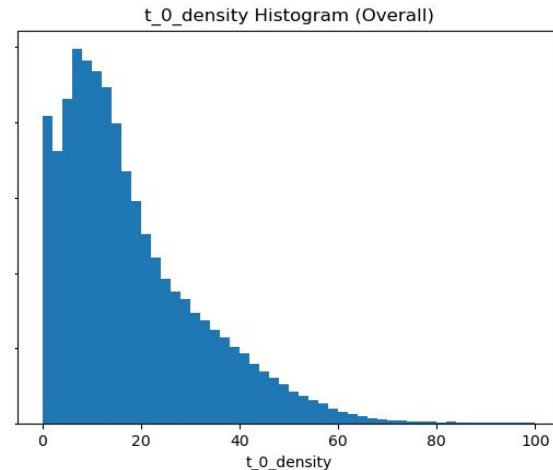
Before Noise



After Noise - All



After Noise - Truncated



Poisson Distribution

Exploratory Modeling Result

In sample data R square, use **fold zero** to calculate out sample data R square

Without Noise

| R Square | Fold 0 | Fold 1 | Fold 2 | Fold 3 | Fold 4 |
|-----------------------|----------------|--------|--------|--------|--------|
| Linear Regression | 0.88 (0.87) | 0.88 | 0.87 | 0.87 | 0.87 |
| Generalized Linear | 0.88 (0.87) | 0.88 | 0.87 | 0.87 | 0.87 |
| Random Forest | 0.82 (0.82) | 0.81 | 0.81 | 0.82 | 0.81 |
| Gradient Boosted Tree | 0.88 (0.87) | 0.87 | 0.88 | 0.88 | 0.87 |

With Noise

| R Square | Fold 0 | Fold 1 | Fold 2 | Fold 3 | Fold 4 |
|-----------------------|----------------|--------|--------|--------|--------|
| Linear Regression | 0.68 (0.83) | 0.88 | 0.80 | 0.85 | 0.87 |
| Generalized Linear | -117 (-77) | -23 | -8 | -76 | -162 |
| Random Forest | 0.11 (0.06) | 0.09 | 0.13 | 0.11 | 0.10 |
| Gradient Boosted Tree | 0.71 (0.55) | 0.68 | 0.61 | 0.62 | 0.63 |

With Noise & Filtering Outlier

| R Square | Fold 0 | Fold 1 | Fold 2 | Fold 3 | Fold 4 |
|-----------------------|----------------|--------|--------|--------|--------|
| Linear Regression | 0.15 (0.12) | 0.61 | 0.25 | 0.36 | 0.38 |
| Generalized Linear | -1.7 (-4.5) | 0.12 | -634 | -15 | -2.6 |
| Random Forest | 0.78 (0.78) | 0.79 | 0.80 | 0.80 | 0.79 |
| Gradient Boosted Tree | 0.84 (0.83) | 0.85 | 0.86 | 0.86 | 0.85 |

With Noise & Add Extra Features

| R Square | Fold 0 | Fold 1 | Fold 2 | Fold 3 | Fold 4 |
|-----------------------|----------------|--------|--------|--------|--------|
| Linear Regression | 0.68 (0.83) | 0.88 | 0.80 | 0.85 | 0.87 |
| Generalized Linear | -119 (-79) | -32 | -5 | -80 | -163 |
| Random Forest | 0.11 (0.07) | 0.10 | 0.13 | 0.12 | 0.11 |
| Gradient Boosted Tree | 0.59 (0.44) | 0.54 | 0.75 | 0.55 | 0.58 |

Shortest Path Algorithm

Model Final Output (linear regression)

- Generate Recommend Path at Week 9 Monday 1PM node o-8
- True Shortest Path
 - 9_30to8_30 8_30to8_31 8_31to8_32 8_32to8_33
8_33to8_34 8_34to8_35 8_35to8_36 8_36to8_37
8_37to8_38 8_38to8_39 8_39to9_39 9_39to9_38
- Predicted Shortest Path
 - 9_30to8_30 8_30to8_31 8_31to8_32 8_32to8_33
8_33to8_34 8_34to8_35 8_35to8_36 8_36to8_37
8_37to8_38 8_38to8_39 8_39to9_39 9_39to9_38
- Source of Error
 - Mainly come from unstable edge (plot on the right)

| edge | MAE | RMSE |
|------------------------|-------------|-------------|
| 6_42to5_42 | 20.46341874 | 33.40086947 |
| Park_54toPark_53 | 16.97218526 | 27.37084394 |
| Park_49toPark_50 | 13.77609672 | 24.13617799 |
| 7_42to8_42 | 13.99103017 | 19.3966747 |
| 6_49to6_50 | 13.06004144 | 18.24747077 |
| 6_34to5_34 | 9.993132689 | 16.80321039 |
| Park_36toPark_35 | 10.96373559 | 16.75342717 |
| 5_52to5_51 | 10.9948562 | 15.8060884 |
| Madison_47toMadison_48 | 9.790172074 | 15.68596128 |
| Madison_57to5_57 | 8.016554576 | 12.39315383 |
| Park_55toPark_54 | 7.663552521 | 11.92563636 |
| Park_38toPark_39 | 9.33225279 | 11.85758805 |
| Park_58toPark_57 | 8.490048711 | 11.19763564 |
| Park_46toPark_47 | 8.33578231 | 11.02605693 |
| Park_47toPark_46 | 7.882120061 | 10.9333112 |
| 5_40to5_39 | 7.775983013 | 10.81476376 |
| Park_42toPark_43 | 7.847863891 | 10.45676273 |
| Park_48toPark_49 | 7.644048104 | 10.02472836 |
| Park_44toPark_45 | 7.315247051 | 9.363016631 |
| Park_50toPark_51 | 7.261492887 | 9.348443243 |
| 8_46to8_47 | 6.782689898 | 8.937917834 |
| 7_37to7_36 | 6.884931109 | 8.80444422 |
| 7_49to7_48 | 6.810883706 | 8.704506702 |
| 8_36to8_37 | 6.769252371 | 8.450968421 |
| Park_52toPark_53 | 6.679018084 | 8.363973546 |
| Park_49toPark_48 | 6.582704206 | 8.33130152 |
| Park_50toPark_49 | 6.097223714 | 7.95653521 |
| 6_36to6_37 | 6.018124317 | 7.942292519 |

Conclusion & Future Steps

All models perform well on no noise data

- Big challenge to make Spark work with a time series problem

After we add noise on data, the story changes...


- Poisson works best for no noise data, not after we add noise and ruin the distribution
- Linear Regression is highly sensitive to the feature distribution and outliers in our noise
- Tree based model is robust for different distribution, but gets penalized by extra useless features
- **It's nearly impossible to find a magic one-size-fits-all model**

Future Steps


- Consider applying different model to different time segments or subgraphs
 - As long as we are clear about the noise source

Project Takeaways


Project Takeaways




1. Container failures and data shuffle may lead to “directory already exists” error when writing data
 - Change write mode to ‘append’
2. Parquet format is better suited for Spark code than CSV format
3. Caching should only be used for high iteration access
4. Data shuffles can make performance suffer
 - Repartition before joins to avoid shuffle
 - Use map instead of repeated joins
 - Data size should be small for broadcast
 - Runtime improved from 3 days to 10 minutes



Project Takeaways



5. Optimizing spark-submit parameters can improve performance
 - Speed up joins and windowing with more executors, cores per executor, and driver/executor memory
 - For caching of large datasets, more driver memory is required
6. Window functions partitionBy and orderBy can be used to compute time-series lag features - with caveats!
7. Cross-validation runs sequentially by default
8. Spark error detection does not catch everything
 - Model trained on infinity values due to floating point error without complaints from Spark
 - Spark ran GraphX code with negative cycles, repeatedly self-healing failed tasks which ultimately could have overloaded the cluster



Demo Instructions

Demo Instructions from README.md

BDAD_Violet_Noise

Team Members

- Ameya - as12366
- Kristin - jl11257
- Emma - yl3750
- Helene - hls327

Running the Demo

Data is designed to be simulated one day at a time, each output from the microsimulation is an ~8GB file. To save time, several pre-processing steps of the data pipeline have already been run so that you can test on a small set of data. Steps completed include generating a new day of data, parsing into parquet format, adding noise, and selecting a subset to test. See the following scripts for details:

```
runscripts/simulateForEval.sh
```

```
runscripts/prepareForEval.sh
```

You are testing on ~2 hours of data, a Monday morning between 7 and 9 AM.

The data is stored in `/user/jl11257/big_data_project/traces/demo/morningsample`

Please Note: The zip file must be unpacked in the top level of your user's `/home/` directory on PEEL. If it is not unpacked at the top level directory, the demo will not work.

To run the demo, simply run `source runscripts/demo.sh` in your home user directory.

If you wish to re-run the demo, you will want to run the following command first.

```
hdfs dfs -rm -r /user/${whoami}/violetnoisestory
```

In the event of folder permission issues please contact hls327@nyu.edu.



Noise and Demo Data



Traffic accidents are simulated to occur between 7 and 10 AM and 2 and 7 PM


- In our Demo data set, we have two accidents in the morning
- We've selected the hours between 7 and 9 AM - which is the noisiest time of day
- Model performance in the demo data is considered to be the worst case scenario
- In addition, data is freshly simulated - it has not been used in any model training

Expected Demo Output

#####

Team Violet Noise
Team Members as12366 yl3750 jl11257 hls327
Traffic Simulation
Demo Run

#####



#####

Checking sufficient permissions...

Input data permissions are OK

Model permissions are OK

Making local hdfs directories for demo output

Running feature calculations for the vehicle classification model, please allow 1-2 mins

Vehicle feature spark log is in vehiclefeaturelog.txt

```
0 0 /user/hls327/violetnoisesummary/vehiclefeatures/ SUCCESS
101.5 K 304.5 K /user/hls327/violetnoisesummary/vehiclefeatures/part-00000-cefb2c02-9ebe-447a-8e29-03e5e6f5fcde-c000.snappy.parquet
100.3 K 300.9 K /user/hls327/violetnoisesummary/vehiclefeatures/part-00001-cefb2c02-9ebe-447a-8e29-03e5e6f5fcde-c000.snappy.parquet
```

Running feature calculations for the edge weight regression model, please allow 1-2 mins

Edge weight feature spark log is in edgefeaturelog.txt

```
0 0 /user/hls327/violetnoisesummary/edgefeatures/ SUCCESS
43.8 K 131.4 K /user/hls327/violetnoisesummary/edgefeatures/part-00000-e654a7b5-d22b-4d28-9aa4-fbac5041637b-c000.snappy.parquet
36.0 K 108.0 K /user/hls327/violetnoisesummary/edgefeatures/part-00001-e654a7b5-d22b-4d28-9aa4-fbac5041637b-c000.snappy.parquet
38.8 K 116.5 K /user/hls327/violetnoisesummary/edgefeatures/part-00002-e654a7b5-d22b-4d28-9aa4-fbac5041637b-c000.snappy.parquet
41.7 K 125.2 K /user/hls327/violetnoisesummary/edgefeatures/part-00003-e654a7b5-d22b-4d28-9aa4-fbac5041637b-c000.snappy.parquet
43.4 K 130.1 K /user/hls327/violetnoisesummary/edgefeatures/part-00004-e654a7b5-d22b-4d28-9aa4-fbac5041637b-c000.snappy.parquet
43.3 K 130.0 K /user/hls327/violetnoisesummary/edgefeatures/part-00005-e654a7b5-d22b-4d28-9aa4-fbac5041637b-c000.snappy.parquet
```

Running car classification predictions on sample data, please allow 1-2 mins

Car prediction spark log is in carclassifylog.txt

There were 48 buses and 12190 cars in the data set provided

The model predicted 147 buses and 12091 cars

Confusion matrix:

```
12088.0 102.0
```

```
3.0 45.0
```

Area under precision-recall curve = 0.29667869149620957

Expected Demo Output


```
Running shortest path forecast on sample data, please allow 2-3 mins
Edge forecast spark log is in edgeforecast.txt
Root Mean Squared Error (RMSE) on out sample test data = 17.77699543550038
Mean squared error (MSE) on out sample test data = 316.02156671380123
Regression through the origin(R2) on out sample test data = -0.5941405438580856
Mean absolute error (MAE) on out sample test data = 9.688900768369908
```

True Shortest Path

```
9_30to8_30 8_30to7_30 7_30to6_30 6_30to5_30 5_30toMadison_30 Madison_30toPark_30 Park_30toLexington_30 Lexington_30toLexington_31 Lexington_31toLexington_32 Lexington_32toLexington_33 Lexington_33toLexington_34 Lexington_34toLexington_35 Lexington_35toLexington_36 Lexington_36toLexington_37 Lexington_37toLexington_38 Lexington_38toLexington_39 Lexington_39toPark_39 Park_39toMadison_39 Madison_39to5_39 5_39to6_39 6_39to7_39 7_39to8_39 8_39to9_39 9_39to9_38
143
```

Predicted Shortest Path

```
9_30to8_30 8_30to8_31 8_31to8_32 8_32to8_33 8_33to8_34 8_34to8_35 8_35to8_36 8_36to8_37 8_37to8_38 8_38to8_39 8_39to9_39 9_39to9_38
201
```



Demo is complete!

#####

Thank You!